

# The building blocks of deep learning models - part 1

A light overview

Filippo Biscarini  
*Senior Scientist*  
*CNR, Milan (Italy)*

Nelson Nazzicari  
*Research fellow*  
*CREA, Lodi (Italy)*



# Deep learning: the building blocks

1. Function approximation
2. The neural network model:
  - a. the “neuron”
  - b. the network
3. Activation functions

today

4. Cost functions
5. Gradient descent (and solvers/optimizers)
6. Forward propagation and the backward propagation algorithm

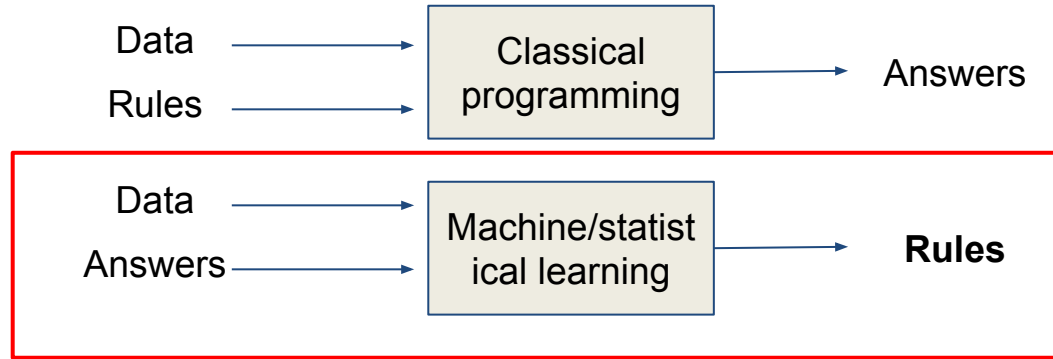
tomorrow



# Function approximation




# Function approximation?



- **unknown function** that maps input data to output results (answers):
  - »  $y = f(x)$
- learn this function → **function approximation**
- $f(x)$  can be **nonlinear** and quite **complex**



# Function approximation




- NNs are good at finding functions that accurately map  $x$  to  $y$
- deep neural networks are powerful **function approximators**
  - »  $y = f(x)$
- complex highly non-linear functions can lead to problems with **generalization**



# Neural network: the “neuron”




# Neural network, the basic unit: the “neuron”





- Mc Culloch & Pitts (1943)
- **perceptron** (“neuron”):
  - dendrites
  - neuron
  - axon



# Neural network, the basic unit: the “neuron”




- Mc Culloch & Pitts (1943)
- **perceptron** (“neuron”):
  - dendrites
  - neuron
  - axon






# Neural network, the basic unit: the “neuron”



- Mc Culloch & Pitts (1943)
- **perceptron** (“neuron”):
  - dendrites
  - neuron
  - axon



# Neural network, the basic unit: the “neuron”



- Mc Culloch & Pitts (1943)
- **perceptron** (“neuron”):
  - dendrites
  - neuron
  - axon
- **learning the weights**


- e.g. linear combination of weights\*features + bias
- fancy way to perform linear regression
- solved through NN rather than OLS or ML



# Anatomy of a neural network




# Fully connected (dense) neural network



- two-layer NN (not strictly “deep”):
  - input layer: [0]
  - hidden layer: [1]
  - output layer: [2]



# Fully connected (dense) neural network




## IMPORTANT!

- each hidden unit takes in input all  $x$  features
- replicates the predictive model as many times as there are units ("neurons")
- if the approximated function is linear regression, each unit will fit a different linear regression model
  - e.g.: 3 units  $\rightarrow$  3 regression models



# Fully connected (dense) neural network



- **two-layer NN** (not strictly “deep”):
  - input layer: [0]
  - hidden layer: [1]
  - output layer: [2]
- all features connected to all “neurons” in the hidden layer
- the NN will decide which variables to use (and how) in each node (by **learning the weights**)




# Activation functions



# Activation functions: what?

"neuron" (unit)




- $g(z)$ : **activation function**
- the unit actually processes both the combination of weights and features and the activation function





# Activation functions: what?


“neuron” (unit)



- $g(z)$ : **activation function**
- the unit actually processes both the combination of weights and features and the activation function
- the output can be i) the final prediction, or ii) the intermediate output of a hidden layer



# Activation functions: when and where?



- **when:** each time a unit is activated: input data (initial features, intermediate output) is processed and output is transferred to the next layer (or final output) through an activation function
- **where:** hidden layers and output layer




# Activation functions: which?

- Identity function
- Logistic function
- Hyperbolic tangent function
- ReLU (Rectified Linear Unit) function
- Softmax function




# Activation functions: identity function




- identity function: a.k.a. **linear activation** function
- returns the value  $z$  that comes from the combination of input features and learned weights
- **never used**, except for the output layer in regression problems




# Activation functions: logistic function



- logistic (**sigmoid**) function
- converts real input in  $[-\infty, +\infty]$  to output in the range  $[0, 1]$




# Activation functions: logistic function



- historically very popular
- now less popular  $\rightarrow$  problems with gradient descent (solution of the model)
- when  $z$  is very large or very small derivatives are close to 0  $\rightarrow$  **slow descent**
- still used for the output layer in binary classification problems




# Activation functions: tanh




- hyperbolic tangent function
- rescaling of the logistic function:

$$\tanh = 2\sigma(2z) - 1 \text{ [proof [here](#)]}$$

- output in **[-1,+1]**, **mean 0**, ~ “centering of the data”



# Activation functions: tanh



- hyperbolic tangent function
- rescaling of the logistic function:


$$\tanh = 2\sigma(2z) - 1 \text{ [proof [here](#)]}$$

- output in **[-1,+1]**, **mean 0**, ~ “centering of the data”
- more efficient learning in the intermediate hidden layers
- still suffers from similar limitations as  $\sigma(z)$  when  $z$  is very large or small







# Activation functions: ReLU




- Rectified Linear Unit function
- derivative is 0 for  $z < 0$ , 1 for  $z > 0$




# Activation functions: ReLU




- most common activation function (default choice in many cases)
- much faster and efficient learning of DL models



# Activation functions: ReLU



$$\max(0, z)$$




## Pros of ReLU activation:

- easy to compute
- sparse representation: many output values will be exactly 0 (unlike sigmoid and tanh, which tend asymptotically to 0)
- avoid vanishing gradients → faster learning (training of multi-layered NNs)


→ ReLU is one of the ingredients that made deep learning possible




# Activation functions: leaky ReLU



- uses a slight slope for  $z < 0$
- can help when there are too many flat neurons (0 slopes, “dying neurons”), e.g.:
  - large negative bias
  - learning rate is too high



# Activation functions: softmax




- returns a probability distribution over the target classes in a **multiclass classification** problem
- $k$  classes
- negative inputs converted to non-negative values (exponential function)
- each input will be in the interval  $[0, 1]$
- same denominator  $\rightarrow$  normalization (sum to 1)

$$\text{softmax}(z_i) = \frac{\exp(x_i)}{\sum_{j=1}^k \exp(x_j)}$$



# Activation functions: softmax



$$\text{softmax}(z_i) = \frac{\exp(x_i)}{\sum_{j=1}^k \exp(x_j)}$$

- Softmax is used in the output layer of multinomial classification problems
- Softmax is differentiable  $\rightarrow$  backpropagation for optimization of the weights (parameters of the deep learning model)



# Activation functions: why not linear?

- the linear (identity) activation function is never used: why?
- has to do with **function approximation**: NNs (deep learning) are excellent at finding complex non-linear relationships in the data (e.g. between features and target variables)
- with the identity activation function, the intermediate output of each layer will just be a linear combination of the input, and so no matter how many hidden layers you have, the final output  $\hat{\mathbf{y}}$  will be a **linear combination** of the initial features  $\mathbf{X}$
- deep learning would then just be a very expensive way of doing linear regression!

$$\begin{cases} y_1 = w_1 x + b_1 \\ y_2 = w_2 y_1 + b_2 \end{cases} \rightarrow y_2 = w_2(w_1 x + b_1) + b_2 = \underbrace{(w_2 w_1)}_{\mathbf{w}'} x + \underbrace{(w_2 b_1 + b_2)}_{\mathbf{b}'}$$

