Supervised learning

Train the learners

Filippo Biscarini
Senior Scientist
CNR, Milan (Italy)

Nelson Nazzicari
Research fellow
CREA, Lodi (Italy)

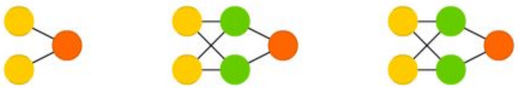

Supervised learning problems

Data

Answers

Statistical learning

Rules

We know the answers (and feed them to the predictive algorithm)

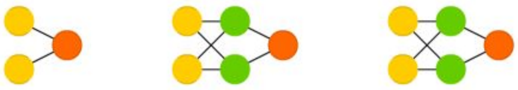

Why supervised?

<u>Training examples</u>

measured variables / features on *n* examples

$$\begin{bmatrix} 1 \\ 0 \\ \vdots \\ 1 \end{bmatrix} = \begin{bmatrix} 0.12 & 1.5 & \dots & 0.9 \\ 2.05 & 0.95 & \dots & 1.1 \\ \vdots & \vdots & \ddots & \vdots \\ 3.5 & 0.88 & \dots & 1.75 \end{bmatrix}$$

labels / target variables (e.g. phenotypes) on n examples

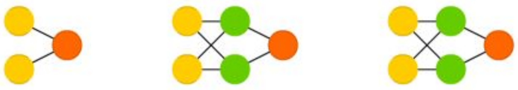

Unsupervised learning

<u>Training examples</u>

measured variables / features
on *n* examples

$$\cancel{\begin{bmatrix} 1 \\ 0 \\ \vdots \\ 1 \end{bmatrix}} = \begin{bmatrix} 0.12 & 1.5 & \dots & 0.9 \\ 2.05 & 0.95 & \dots & 1.1 \\ \vdots & \vdots & \ddots & \vdots \\ 3.5 & 0.88 & \dots & 1.75 \end{bmatrix}$$

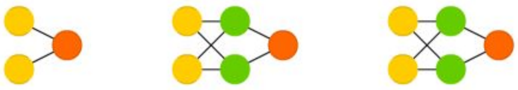

Unsupervised example: K-means clustering

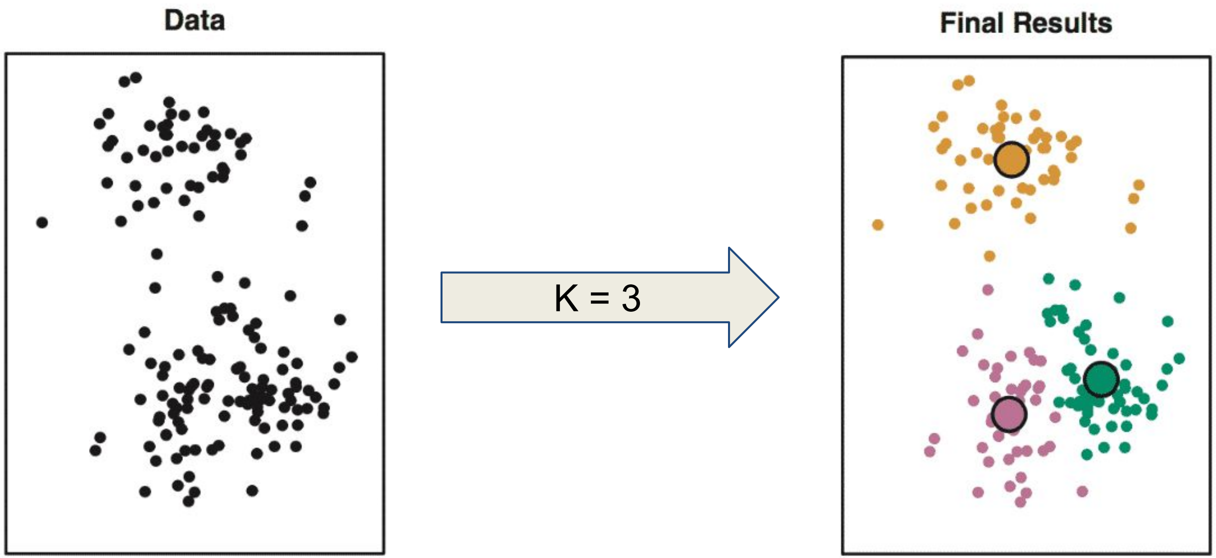

source:
https://www.iotforall.com/machine-learning-crash-course-unsupervised-learning

Regression and classification

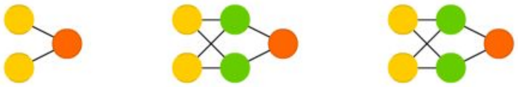

Supervised learning problems

- Regression (**predictive**) problems
- Classification (**predictive**) problems

Supervised learning problems

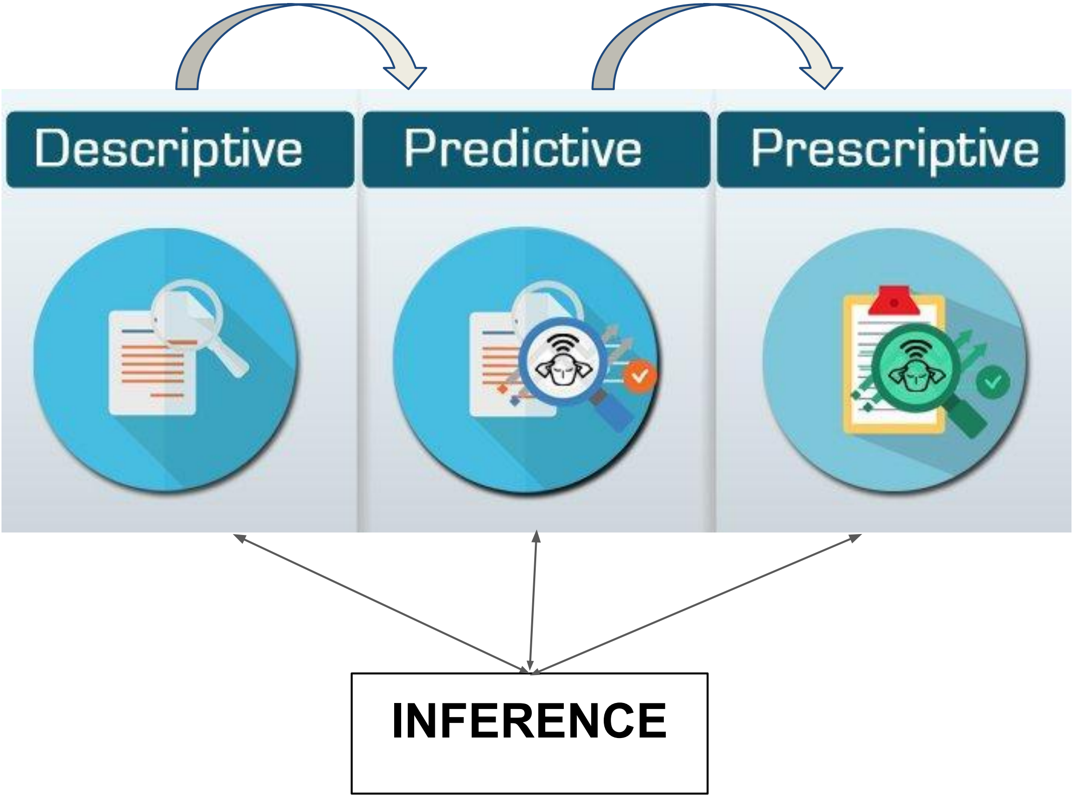

- Know the past
- Predict the future
- Act consequently

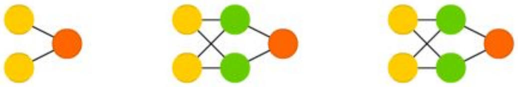

Supervised learning problems

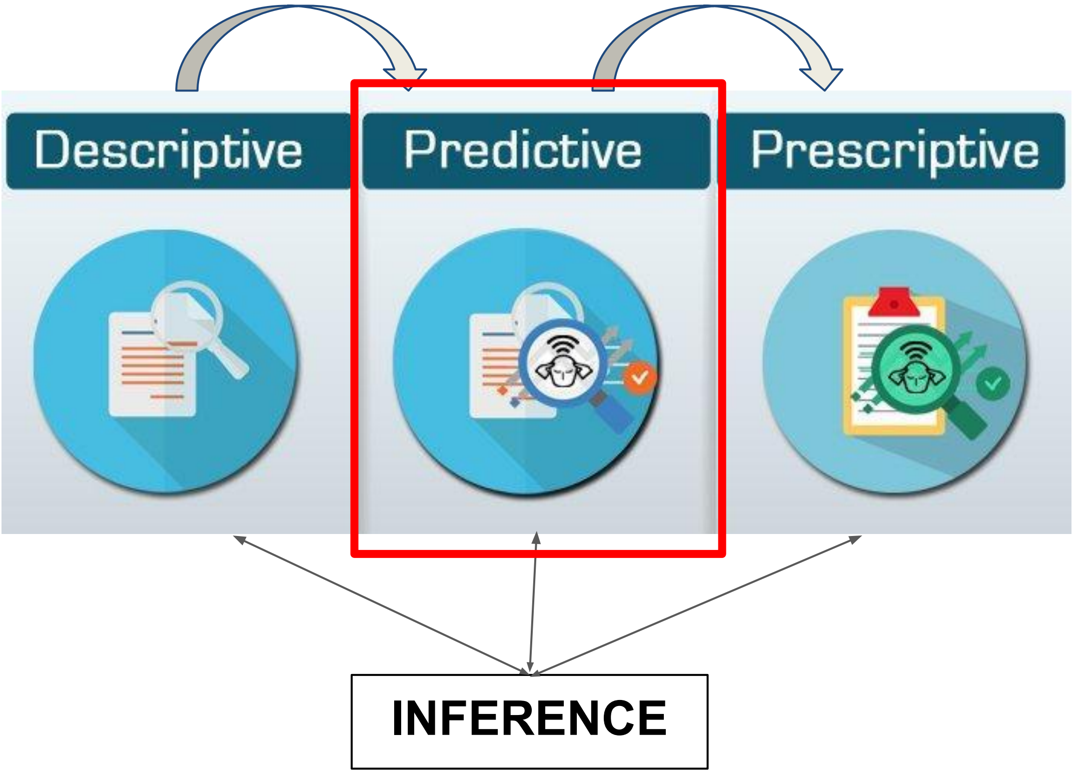

- Know the past
- Predict the future
- Act consequently

- A catch-all term
- Can be confusing

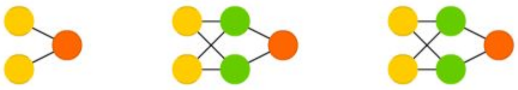

Inference vs Prediction

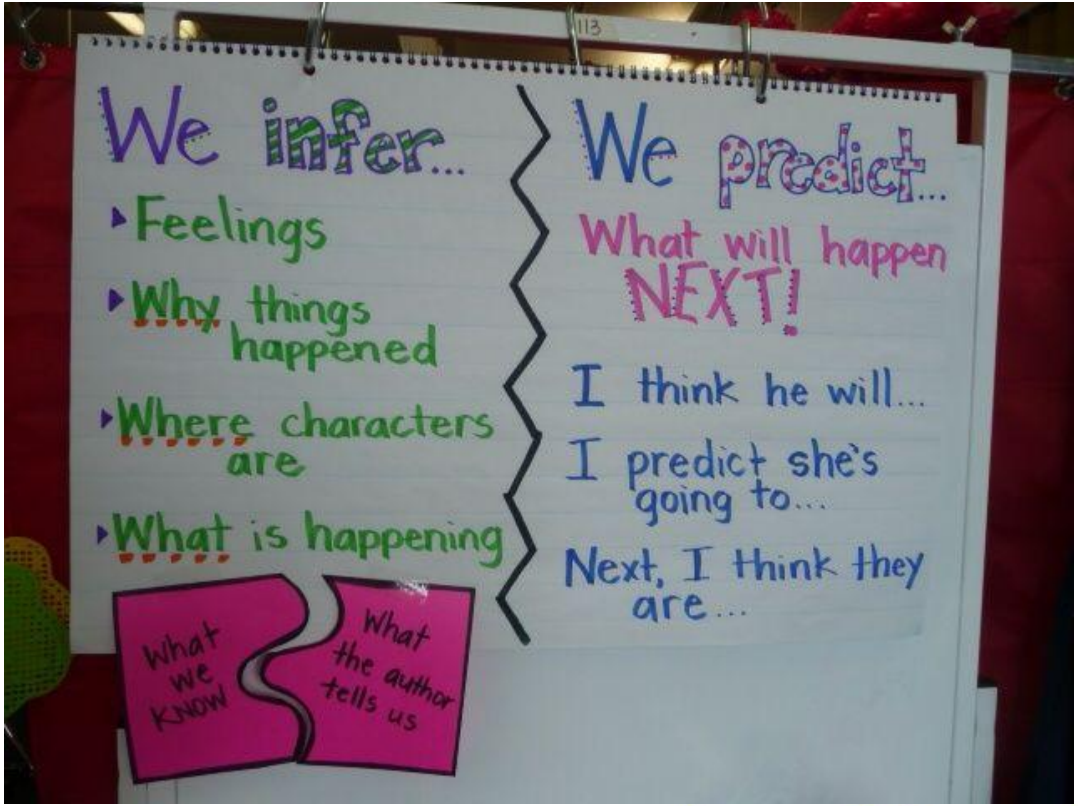

- different statistical problems
- different objectives, different rules … different ballparks
- inference is in general more difficult than prediction

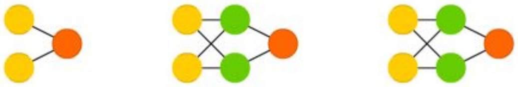

Supervised learning problems

- Regression (**predictive**) problems
- Classification (**predictive**) problems

Predictive machines!

- Classifiers
- Predictors

source:
https://blog.bigml.com/2013/03/12/machine-learning-from-streaming-data-two-problems-two-solutions-two-concerns-and-two-lessons/

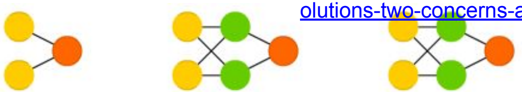

Regression problems

- the response variable **y** is **quantitative**
- e.g.: *height*, *weight*, *yield* (milk, crops), *blood sugar concentration*
- **y** = **target** (dependent) variable (a.k.a. response, objective variable)
- **X** = matrix of **features** (continuous, categorical)
- **predictor**: y = f(x) = **P(X)** ← [predictive machine]

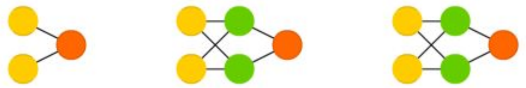

Regression problems - simple regression

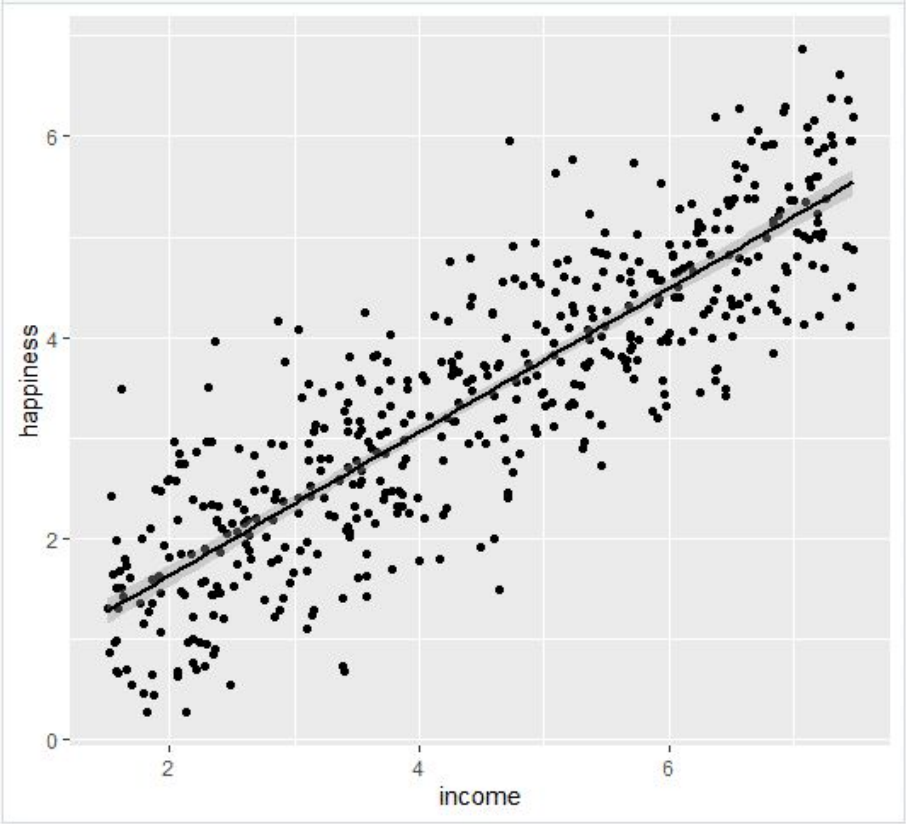

happiness = (intercept) + beta*income

or

income = (intercept) + beta*happiness

Source: https://www.scribbr.com/statistics/linear-regression-in-r/

Regression problems - multiple regression

$R^2 = 0.79$

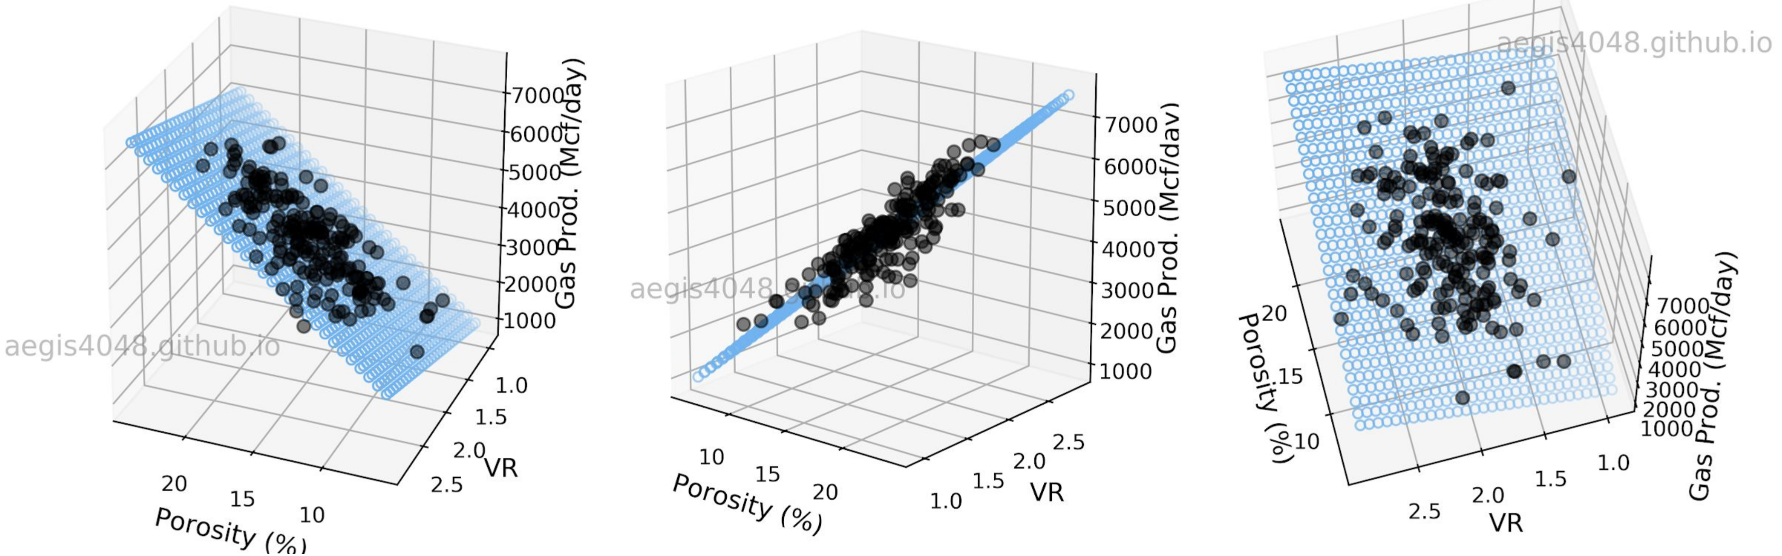

Source: https://aegis4048.github.io/mutiple_linear_regression_and_visualization_in_python

Multiple linear regression

$$y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \ldots + \beta_p X_p + \epsilon$$

- y: target variable
- **β**'s: model coefficients
- X's: features (predictors, independent variables, factors)

Multiple linear regression

$$\mathbf{y} = \beta\mathbf{X} + \mathbf{e}$$

- matrix (compact) notation
- vectors of observations (**y**), coefficients (**β**) and residuals (**e**)
- matrix of features (**X**)

Multiple linear regression

$$\mathbf{y} = \beta\mathbf{X} + \mathbf{e}$$

estimation of coefficients

$$\hat{\mathbf{y}} = \hat{\beta}\mathbf{X}$$

→ **predictions!**

- matrix (compact) notation
- vectors of observations (**y**), coefficients (**β**) and residuals (**e**)
- matrix of features (**X**)

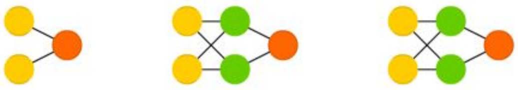

Predictions

$$\hat{y} = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \ldots + \beta_p X_p$$

with the estimated coefficients $\boldsymbol{\beta}$ and the feature values **X** we obtain the predicted values $\hat{y}$

Predictions

$$\hat{y} = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \ldots + \beta_p X_p$$

with the estimated coefficients $\boldsymbol{\beta}$ and the feature values **X** we obtain the predicted values $\hat{y}$

→ how do we obtain the model coefficients β?

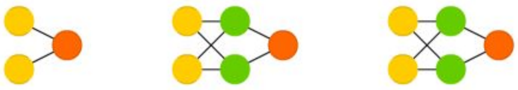

Supervised learning recap

- **deep learning** is one of many methods that can be used to solve supervised learning problems
- deep learning is mainly used in **predictive problems**
- (could be used though also for inferential problems and for unsupervised learning)
- we'll see later how to use deep learning to solve linear regression